

CIVIC BRIDGE AI: FROM HATE SPEECH DETECTION TO TRANSFORMATIVE DEMOCRATIC EDUCATION IN DIVIDED SOCIETIES

Challenge selected: Think out of the box

Working MVP: <https://civic-bridge-ai-r3qs2dc4cngr9oftowff5.streamlit.app/>

Team: Prof. Charis Psaltis, Dr Demetris Paschalides, Dr Anna Zapiti, Mr Maximos Psaltis

Human rights advisor: Prof. Aristotelis Constantinides, newly elected Chairperson of the Management Board of the European Union Agency for Fundamental Rights (FRA)

Educational Advisor: Dr Chara Makriyiani, EUROCLIO Ambassador, School Principal, Chair of the Scientific Advisory Council (SAC) of the CoE Observatory on History Teaching in Europe (OHE) (2021-2025)

Media Advisor: Voula Antoniou, Media Producer & Director

Problem

Hate speech not only offends individuals, but especially in divided societies, it hardens the symbolic boundaries of democratic dialogue, turns social groups into enemies, transforms legitimate fears into collective blame, and blocks citizens from engaging with alternative perspectives. Current approaches often focus on detecting, reporting, or removing harmful content. These responses are necessary in serious cases. Nevertheless, they are not always sufficient to address the social and educational mechanisms that reproduce hate, such as one-sided narratives, the absence of multiperspectivity, lack of critical media literacy, and weak skills for democratic disagreement. The 2026 Hackathon of the Council of Europe (CoE) asks teams to move beyond traditional moderation and design solutions that prevent, contextualise, counter, or transform hate speech. *Civic Bridge AI* responds directly to this challenge. It is a privacy-conscious, human-in-the-loop web prototype that turns hate speech detection in social media and news portals into democratic learning and response.

Solution

Civic Bridge AI is a theory and CoE policy based App that analyses public or anonymised text and guides a variety of possible users through five steps: **1. Detect → 2. Interpret → 3. Contextualise → 4. Educate → 5. Respond**

The MVP allows users to upload a CSV of public posts or anonymised examples - also expandable to be fed API connections to media platforms. It then classifies each text into a graduated severity model inspired by Council of Europe Recommendation CM/Rec(2022)16:

1. ordinary political expression
2. offensive or harmful expression
3. potential hate speech for human review
4. high-severity incitement/dehumanisation risk

The system does not make legal determinations. It produces risk-support signals for trained human review, with strong warnings against automated censorship or punitive action. The innovation lies in what happens *after* detection. *Civic Bridge AI* identifies narrative frames, informed by Genetic Social Psychology (an interdisciplinary framework of social developmental psychology and history teaching-CUP open access: <https://www.cambridge.org/core/books/conflict-and-change/E4AF5714AD8DC0BDB97753965EF5924B>) and dialogical psychology. Moreover, it detects semantic barriers, that is, discursive mechanisms that prevent engagement with alternative views, such as rigid opposition, stigma, distrust, transfer of meaning, prohibited thoughts, bracketing, and undermining the motive of the author. It then proposes semantic promoters - dialogical strategies that help users reopen democratic communication, including recognising ingroup bias, contextualisation, corroboration, outgroup empathy, superordinate identity, and condemnation of harm regardless of perpetrator based on social psychological theories related to intergroup contact and prejudice reduction.

The platform includes a dedicated education module with age-differentiated guidance for:

ages 6-9	ages 10-13	ages 14-17	ages 18+
----------	------------	------------	----------

For educators and parents, it provides developmental notes, RFCD-linked competences, activity suggestions, and “what to watch for” guidance. For pupils and students, it offers age-appropriate prompts to examine harmful statements, identify propaganda mechanisms, and produce bridge formulations or counter-narratives.

How AI Is Incorporated:

Civic Bridge AI is a large-language-model-orchestrated pipeline with a human in the loop, not a standalone moderation classifier. The model performs structured analysis at each of the five stages, with a trained human retaining every decision that carries consequences. Rather than a binary “hate / not hate” label, the system prompts the model with the graduated severity criteria of CM/Rec(2022)16, returning a tier placement and a short, criteria-anchored rationale. Beyond severity, the model identifies the discursive mechanisms (semantic barriers) that shut down dialogue and, conditioned on these and the selected age band, proposes promoters and candidate bridge formulations for a human to adapt. In the current MVP this analysis is performed through prompt-structured frame detection: the roadmap replaces it with learned clustering validated against expert human coding. A prompt-driven LLM is preferred because it generalises across 46 member states and multiple languages without per-country training data and produces interpretable rationales that a non-technical reviewer can audit. The model narrows attention and explains its reasoning, and a trained human decides.

Country and Language Context

The MVP includes a country and language selector for all 46 Council of Europe member states. Cyprus is developed as the first case study, with Greek, Turkish, and English context modules. This is deliberate: Cyprus is a protracted conflict setting inside the Council of Europe area, where intercommunal counter-narratives, anti-migrant discourse, anti-Muslim rhetoric, anti-Black racism, anti-LGBTI stigma, and gendered attacks can intersect.

The Cyprus module draws on ECRI's Sixth Report on Cyprus, which identifies the absence of comprehensive hate-speech monitoring, sporadic counter-speech by public figures, and weak implementation of anti-racist education policies. It also incorporates youth-facing resources such as the Cyprus Police Charter of Rights for Adolescents, including online safety and help-seeking guidance. The Turkish-language Cyprus module is carefully framed for Turkish-speaking users in Cyprus, including Turkish Cypriots and users living in the northern part of Cyprus. It avoids terminology that could imply recognition of any separate state entity and keeps the country context as Cyprus, using language consistent with European institutional practice.

User Guidance

A built-in *Guidance Bot* helps users decide what to do depending on their role and need. It asks whether the user is an educator, student, civil society actor, journalist, public institution, or targeted person/witness. It then provides next steps for situations such as seeing harmful content online, creating counter-speech, using an example in education, understanding the local context, or seeking help/reporting. For urgent situations, it prioritises emergency help and evidence preservation rather than public engagement.

Privacy and Human Rights Safeguards

Civic Bridge AI is designed to work without usernames, account IDs, profile links, exact locations, or private messages. It analyses discourse patterns, not user identity. It encourages anonymised or synthetic data for educational use and preserves only minimal public-text metadata where needed. This responds to the privacy-preserving hate speech detection challenge while staying aligned with freedom of expression, proportionality, dignity, non-discrimination, and human review.

What the system does not do: *Civic Bridge AI* does not make legal determinations about whether speech is unlawful; does not automatically remove, hide, or report content; does not take any action against a user; and does not operate without trained human review. It produces explained risk-support signals and educational material for humans to evaluate. It analyses discourse, not identity.

Implementation at the Hackathon

During the Hackathon we will improve the existing *Streamlit MVP* by:

1. strengthening the LLM coding prompt with CM/Rec(2022)16 severity criteria;
2. improving the Cyprus country/language modules;
3. refining the semantic barrier/promoter detection;
4. adding a more polished educator/student workflow;
5. preparing a five-minute demonstration using anonymised Cyprus examples;
6. documenting privacy, human-rights, and non-recognition safeguards.

The longer-term roadmap includes multilingual testing, learned semantic clustering (current version uses prompt-structured frame detection), and validation of the AI signals against expert human coding

Team: Our team combines democracy and policy expertise, technical capacity, education design, and a youth perspective. Psaltis brings expertise in social and developmental psychology, intergroup relations, history education, conflict transformation, and democratic competences. Paschalides brings technical expertise in AI/NLP and prior experience from MANDOLA, a European project on monitoring and detecting online hate speech. Zapiti contributes educational and psychological expertise for age-appropriate interventions, with projects and educational material on human rights, democracy, and racism. Maximos Psaltis, an aspiring electric engineer and a 17-year-old secondary school student, provides an authentic youth-user perspective for the student-facing module. Constantinides advises on human rights, international law, freedom of expression, and fundamental rights safeguards. Makriyianni and Antoniou are our advisors for education and media, respectively.

Impact: *Civic Bridge AI* is a multiuser multitool that helps make democratic security practical and strives to make an impact on multiple fronts. It supports educators, young people, civil society, journalists, and institutions to move from "How online hate speech can be detected?" to "How can this harmful narrative be understood, challenged, and transformed?" It contributes to the New Democratic Pact for Europe by combining technology, democratic education, human rights safeguards, and local context. Its purpose is not to automate judgement, but to help people recognise hate, protect dignity, and rebuild the conditions for democratic dialogue.

Sources for Alignment

- Council of Europe Hackathon page: <https://www.coe.int/en/web/new-democratic-pact-for-europe/2026-hackathon>
- Council of Europe Recommendation CM/Rec(2022)16 on combating hate speech: https://search.coe.int/cm/pages/result_details.aspx?objectid=0900001680a67955
- ECRI Cyprus country page: <https://www.coe.int/en/web/european-commission-against-racism-and-intolerance/cyprus>